

# Go`ing, go`ing.. gone!

Some bits on the Go runtime

June 5th, 2014 for the GoTO meetup group

Peter Kieltyka / [peter@pressly.com](mailto:peter@pressly.com)

# Personal exploration


- Writing backend software for the last 10 years
- Started with writing small UNIX network apps in C
- Wrote web apps with PHP for a few years, discovered and fell in <3 with Ruby, tried to make it fast for years after that
- Tried Node, Scala..
- Found Go

Go's goal: to make writing  
concurrent server software  
easy, safe and fun

# Go's concurrency model

- Inspired by Tony Hoare's CSP (Communicating Sequential Processes)
- Message passing vs. shared state
- The concurrency primitives offered by Go:
  - goroutines & channels
- No “thread” api for application developers; that's all under the hood (runtime)

# The two important pieces for fast network applications



# Multithreading models..

**1:1**

**N:1**

**M:N**

Kernel-level threading

User-level threading

Hybrid threading

OS threads

Green threads

OS threads + green  
threads

OS Scheduler

App runtime  
scheduler

OS / App runtime  
scheduler


Threads are  
expensive (initial size,  
1 to 8 mb per thread)

Fast context  
switching, but, what  
about the other  
cores?


Difficult to implement,  
balance of language  
syntax / runtime

# Go's Scheduler

The scheduler's job is to distribute ready-to-run goroutines over worker threads. The main pieces:



Machine



Processor



Goroutine




# Goroutine

```
struct G {  
    int64      goid;          // Go id  
    uintptr    stackbase;     // The stack  
    uintptr    stack0;        // Stack pointer  
    uintptr    stacksize;     // Stack size  
    int16      status;        // Current state in runtime  
    ...  
} // src/pkg/runtime/runtime.h
```

- The `go fn()` we all know and love; spawn a goroutine to later execute 'fn'. Puts a G on the local run queue.
- 8KB per goroutine on initial size (Go 1.2 and 1.3)
- Stack grows and shrinks as necessary
- `main()` is also a G
- Status (runtime.h): Idle, Runnable, Running, Syscall, Waiting, Dead
- Executes on an **M**achine in the context of a **P**rocessor





# Machine

```
struct M {  
    int32  id;           // Machine id  
    bool   spinning;    // M is out of work and is actively looking for work  
    bool   blocked;     // M is blocked  
    int32  gcing;       // M is garbage collecting  
    G*     curg;        // Current running goroutine  
    P*     p;           // Attached P for executing Go code (nil if not executing Go code)  
    ...  
} // src/pkg/runtime/runtime.h
```

- Managed OS Thread for executing Go code
- Start/stop new M's (threads), as necessary, managed by the runtime not the developer
- A thread cache for efficiency
- Special rules for when blocking on a Syscall or executing a CGO function

# P

# Processor

```
struct P {  
    int32      id;           // Processor id  
    uint32     status;       // Current state: Idle, Running, Syscall, Gcstop, Dead  
    M*         m;           // Current associated Machine  
    G*         runq[256];    // Goroutines queue  
    ...  
} // src/pkg/runtime/runtime.h
```

- Exactly `GOMAXPROCS` P's
- P schedules Goroutines to run on an attached M in FIFO order
- Context for scheduling G's on M's, must attach to an M to execute a new G (code)
- P has local run queue of Goroutines, as well periodically checks a global queue of G's
- and, when a P's go run queue is empty, it will steal work from other busy P's

# Peak under the hood

ex.go

```
package main

import (
    "fmt"
    "time"
)

func main() {
    fmt.Println("hey hey, my main G")
    go myFn("hey, wait up")
    time.Sleep(10e6) // give it a sec..
}

func myFn(s string) {
    fmt.Println(s)
}
```

./ex (from proc.c)

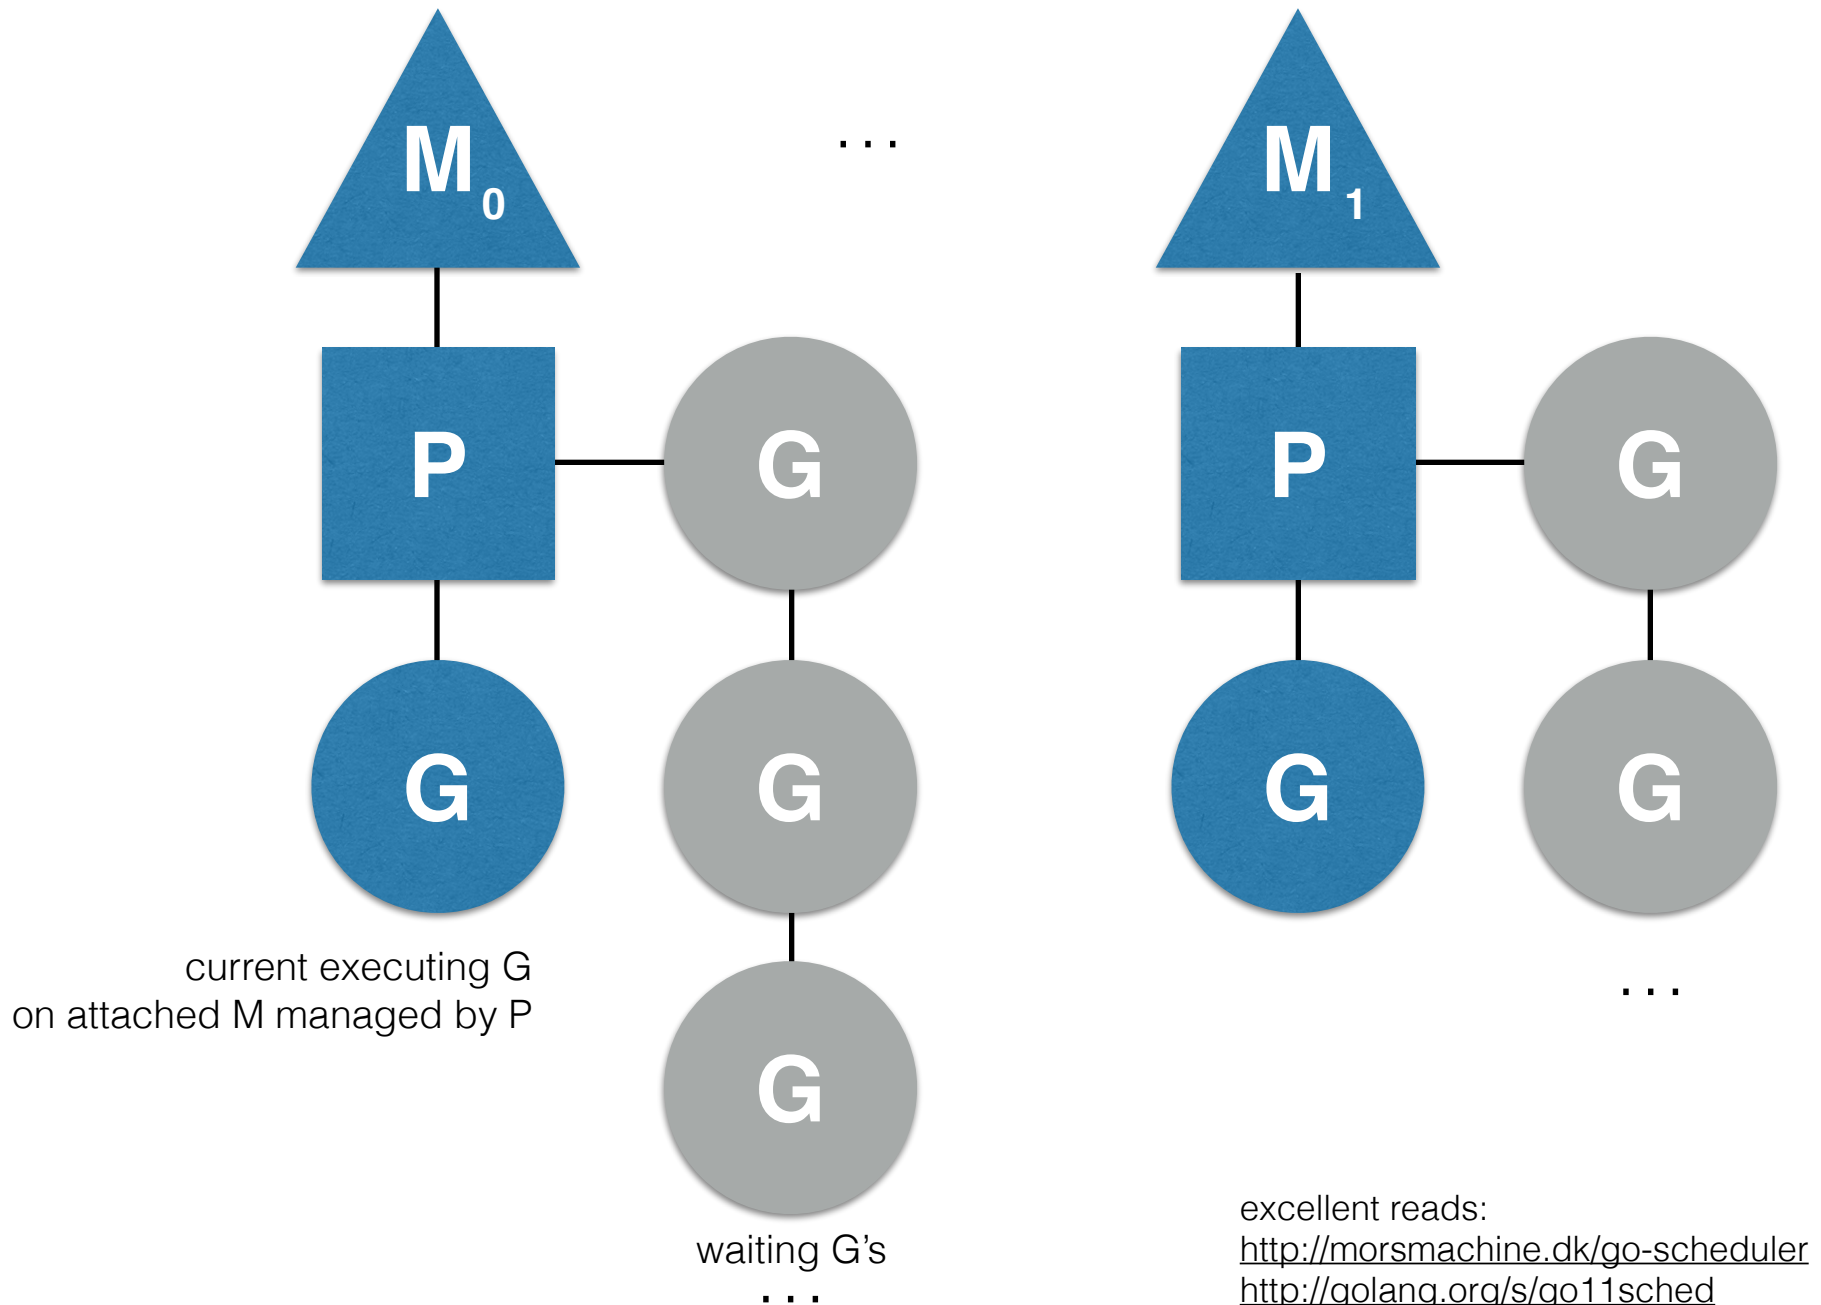
Bootstrap:

- runtime·args()
- runtime·osinit()
  - Set runtime·ncpu
- runtime·schedinit()
  - Set max M count (10k default)
  - Set args / env vars
  - Init first P
- runtime·newproc(&mainFn())
  - Make a new G and queue it
  - Effectively calling `go main()`
- runtime·mstart()
  - Init a new M
  - Find a new G to run
  - Execute the G

Exec main G:

- fmt.Println(..)
- runtime·newproc(&myFn())
  - Make a new G for call to myFn and queue it
- time.Sleep(..)


# Scheduler visual; GOMAXPROCS=2



# Here's what's really interesting (IMO)..

- So, the scheduler abstracts threads with Goroutines, and manages their execution across multiple cores
- On a network operation, a Goroutine will block, and be put back on the run queue until its ready
- Blocking..? how is it so efficient and fast?
- Go actually uses an evented IO under the hood whenever performing network operations.. just like Nodejs, Scala, etc.

# A real world example of evented IO kicking butt



\*Google it

# Netpoller

- Network IO interface to kqueue, epoll, IOCompletionPort
- Kernel-level stateful file descriptor monitoring and event notification of readiness and completion
- As opposed to using select()/poll() which have to check each fd each iteration
  - With thousands of operations, this becomes expensive

# Example

```
//...package, imports..  
func main() {  
    go func() {  
        for {  
            time.Sleep(10 * time.Millisecond)  
            fmt.Println("tick..")  
        }  
    }()  
  
    var wg sync.WaitGroup  
    wg.Add(1)  
    go func() {  
        resp, _ := http.Get("http://www.google.ca/")  
        defer resp.Body.Close()  
        fmt.Println(resp.Status)  
        wg.Done()  
    }()  
  
    wg.Wait()  
}
```

Block G on network io operations;  
put the G back on the local runq and set  
g->status = waiting



The beauty of Go to me is in  
the design of language  
(syntax), the threading model  
and evented network IO

Writing speedy servers has  
never been easier.



Thanks!

[peter@pressly.com](mailto:peter@pressly.com) | @peterk